LARGE LANGUAGE/GENERATIVE AI

Practical AI Use Guide

Empowering Non-Technical Teams for Everyday Excellence

Prepared By: Madhumitha D/O Suresh Kumar

Date: July 8, 2026

Table of Contents

1. Quick Start: What Makes a Strong Prompt?

Prompting is the act of guiding an Artificial Intelligence to get the specific output you need. Think of the AI as a highly competent, newly hired intern: it possesses immense knowledge but lacks context regarding your internal goals, style, and boundaries. To get the best results, you must avoid short, vague commands and instead build structural, comprehensive prompts.

Use the following framework to assemble strong, high-yield prompts every time:

> 🎯 **Task:** Clearly describe what you want the AI to do. Use explicit action verbs (e.g., 'Summarize', 'Draft', 'Analyze', 'Categorize').

> 📖 **Context:** Provide crucial background information. Explain the situation, why you are doing this, and any history the AI needs to know.

> 👥 **Audience:** Specify exactly who will read or use the output (e.g., executive board, angry customer, new hires, non-technical vendors) so the tone adjusts appropriately.

> 📊 **Format:** Define what the visual or structural layout should look like (e.g., Markdown table, bulleted list, email template, 3-sentence summary).

> ⛔ **Constraints:** Set clear boundaries, rules, and limitations (e.g., 'Keep it under 200 words', 'Do not mention pricing details', 'Avoid corporate jargon').

> 🔬 **Evidence:** Provide external data, reference text, or exact metrics for the AI to lean on rather than letting it generate facts freely.

> 🔍 **Review:** Instruct the AI to evaluate its own output or guide yourself on how to verify accuracy before deploying it.

2. Weak Prompt / Strong Prompt Comparison

To see how the framework functions in professional, real-world workplace scenarios, explore the structural upgrades below:

Scenario A: Customer Service (Handling a Refund Request Delay)

Weak Prompt:
❌ *Write an email to a customer whose refund is late.*

Strong Revised Prompt:
🚀 Task: Draft a polite email explaining a 4-day delay in processing a refund.
Context: Our billing system underwent an unexpected update yesterday, pausing automated transactions.
Audience: A loyal premium customer who has contacted us twice about their missing $150 credit.
Format: Professional 3-paragraph email template with placeholders for [Customer Name] and [Order ID].
Constraints: Maintain an empathetic, professional tone. Do not offer discount codes, but reassure them the issue is prioritized.

💡 The weak prompt leaves the AI guessing the tone, customer profile, and specific reason for the late refund. The strong version prevents generic text by establishing an empathetic tone for a premium customer, giving an accurate reason, and controlling the formatting precisely.

Scenario B: Human Resources (Onboarding Guide Summary)

Weak Prompt:
❌ *Summarize this onboarding document for new employees.*

Strong Revised Prompt:
🚀 Task: Condense our 15-page internal onboarding guide into a quick-reference summary.
Audience: Incoming non-technical associates on their very first day of work.
Format: A bulleted list divided into three clear categories: 'Day 1 Tasks', 'Week 1 Goals', and 'Key Contacts'.
Constraints: Highlight only action items. Omit complex compliance codes or legal jargon. Keep the total length under one page.

💡 A generic summary often leaves out critical timelines and mixes crucial tasks with legal fine print. By structuring the format into thematic blocks and demanding the exclusion of legal jargon, the output becomes immediately actionable for a day-one employee.

Scenario C: Project Planning (Kickoff Meeting Agenda)

Weak Prompt:
❌ *Make a meeting agenda for our new project launch.*

Strong Revised Prompt:
🚀 Task: Create a highly structured 45-minute project kickoff meeting agenda.
Context: We are launching a digital client portal next month. Cross-functional teams (Product, QA, Marketing) will be present.
Format: A clean timeline table showing Time Allocated, Topic, Speaker Role, and Desired Outcome.
Constraints: Dedicate exactly 10 minutes to risks/blockers. Ensure the meeting concludes with a clear 5-minute next-steps confirmation section.

💡 Instead of a generic bulleted list of topics, the strong prompt forces time boundaries, assigns roles, demands a focus on blockers, and produces an execution-oriented table tailored to cross-functional alignment.

3. Responsible AI Use Checklist

Before copying, pasting, or sharing any AI-generated text, run it through this operational quality control checklist:

[ ] 🔒 Privacy & Data Protection
Ensure absolutely zero customer PII (Personally Identifiable Information), vendor contracts, sensitive financials, or proprietary source code were pasted into public AI models.

[ ] 📅 Source Quality & Currency
Does this task require real-time market data or highly recent regulatory changes? Ensure the model isn't relying on outdated pre-training knowledge.

[ ] ✅ Accuracy & Fact-Checking
Did you manually verify names, dates, calculations, and empirical assertions against an official internal or external truth source? (AI models can hallucinate plausible-sounding errors).

[ ] ⚖️ Bias & Fairness Review
Read the output critically. Does the text make unsupported assumptions or include loaded words that could treat a demographic, group, or employee unfairly?

[ ] 📢 Transparency & Disclosure
Does company policy or client consensus require a disclosure note stating: 'Automated tools contributed to the formulation of this text'?

[ ] 🧠 Human Judgment & Ownership
Is there a human manager or domain expert assigned to take final accountability for this document before it goes live to clients or leadership?

4. Tokenization and Context Windows

What is Tokenization?

AI does not read words the way human beings do. Instead, it breaks down paragraphs into small linguistic fragments called tokens. A token can be a single character, a syllable, a whole word, or part of a complex term. On average, 1 token is roughly equal to 4 characters or 0.75 words.

Visual Token Breakdown Example:

`Prompt` `ing` `helps` `work` `flows.`

(The phrase 'Prompting helps workflows.' is divided into distinct structural tokens processed simultaneously by the engine.)

Why Long Conversations Affect Performance

Every AI model possesses a **Context Window**, which is the total capacity of tokens it can read and remember during a single session. Crucially, this container is shared: both your prompts (inputs) and the AI's responses (outputs) drain the exact same pool of tokens simultaneously. Think of it like your own short-term memory even. When you provide overly long documents, complex prompts, or engage in a continuous, multi-turn conversation, this memory gets full, leaving little room for reasoning and response. It can lose track of your earlier rules, miss critical parameters, or start hallucinating information because its processing focus is overextended resulting in degrading performance.

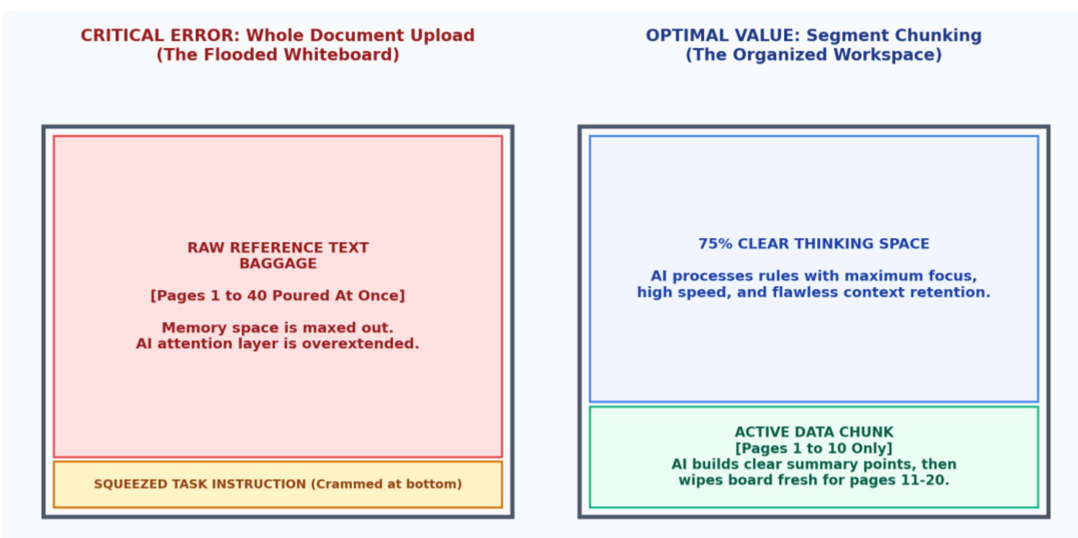

Figure 4.1: Visual layout contrasting Whole Document Upload with the Segment Chunking approach on an Office Whiteboard.

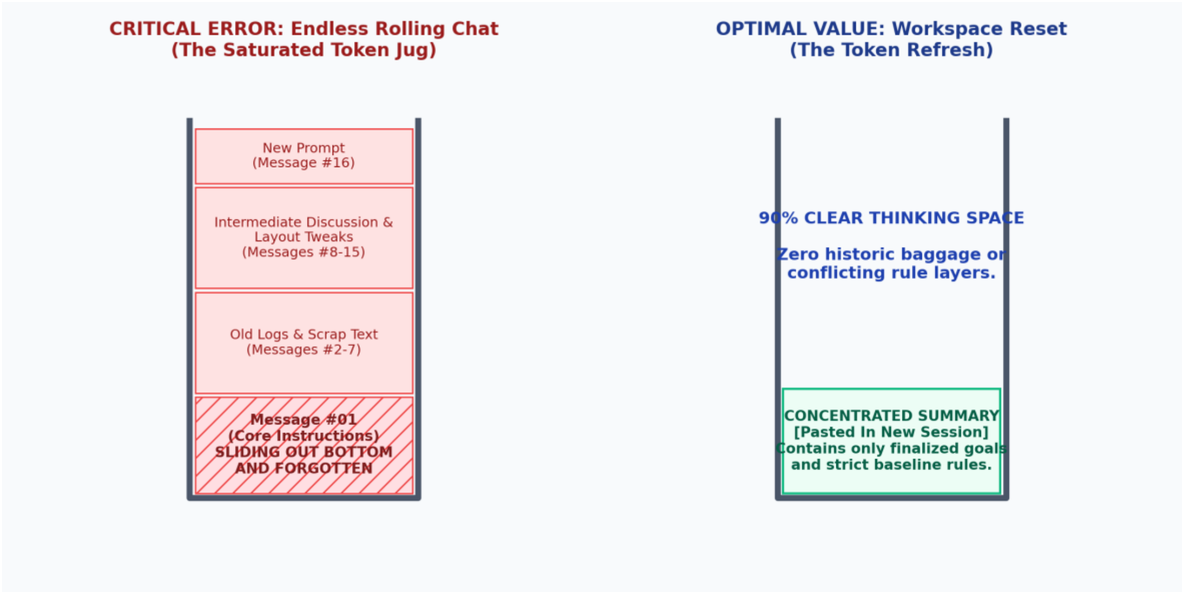

Figure 4.2: Visual illustration of how earlier context gets lost when tokens are used up in an extended chat.

🛠 Concrete Operational Thresholds & Actionable Solutions:

> 🧹 **Prompt Clearance (Trigger: Material with >15% boilerplate):** Manually strip legal signatures, disclaimers, email chains, and decorative characters before pasting. This preserves high-value memory tokens for critical analytical tasks.

> 🧩 **Segment Chunking (Trigger: >10 pages / 4k words):** Instead of dropping a full document, feed it into 10-page sections. Prompt: 'Read this first section, list the top 3 themes, say "READY FOR PART 2", and stop. Wait for me to provide the next section before running final analytics.'

> 📝 **Interim Summarization (Trigger: >15 chat turns):** When chat depth gets heavy, prevent amnesia by prompting: 'Summarize our agreed constraints and goals in a clear list.' Copy that summary, open a fresh New Chat window, paste it, and proceed with a clean whiteboard memory.

> 🎯 **Section-by-Section (Trigger: Long deliveries >4 pages):** If the output is expected to be large, force the AI to perform heavy structural layout operations systematically. Demand Section 1 first, refine it, approve it, and only then issue instructions to proceed with Section 2.

5. Prompting for Grounded Answers

To avoid 'hallucinations' (instances where AI invents fictional facts or mixes up policies), you should restrict the model's operating environment to a closed loop. This technique is known as Grounded Answer Prompting or standard manual context anchoring. By providing the source material explicitly and forbidding the AI from using its general knowledge base, you protect operational integrity.

An Example Grounded Blueprint:

> You are an internal HR assistant operating strictly under corporate compliance guidelines.
>
> TASK: Answer the employee question below using ONLY the provided Policy Text block.
>
> CONSTRAINTS / RULES:
> 1. Do not use external internet knowledge or make inferences.
> 2. If the answer cannot be found verbatim within the provided Policy Text, state exactly: "The provided policy does not include that information."
>
> [START OF POLICY TEXT]
> Employees are eligible for full remote work stipends after completing 90 consecutive days of full-time employment. The stipend covers up to $50 per month for verified broadband internet bills. It does not cover mobile phone data plans, office chairs, or hardware accessories.
> [END OF POLICY TEXT]
>
> EMPLOYEE QUESTION: Can I use my monthly stipend to buy an ergonomic office chair?
> YOUR ANCHORED ANSWER: